

## 1 Question 1

Compute (i) forward, (ii) backward, and (iii) central gradient image. Using the gradient compute the gradient magnitude image and gradient direction image. Analyze the difference you observed in the magnitude and direction images.

### Answer 1

For the forward difference, the gradient in the  $x$  and  $y$  directions is computed as follows:

$$\begin{aligned} G_x(i, j) &= I(i, j + 1) - I(i, j) \\ G_y(i, j) &= I(i + 1, j) - I(i, j) \end{aligned}$$

For the backward difference, the gradient in the  $x$  and  $y$  directions is computed as:

$$\begin{aligned} G_x(i, j) &= I(i, j) - I(i, j - 1) \\ G_y(i, j) &= I(i, j) - I(i - 1, j) \end{aligned}$$

The central difference method averages the forward and backward differences:

$$\begin{aligned} G_x(i, j) &= \frac{I(i, j + 1) - I(i, j - 1)}{2} \\ G_y(i, j) &= \frac{I(i + 1, j) - I(i - 1, j)}{2} \end{aligned}$$

#### 1.1 Forward Difference

As mentioned above we now develop the python code to compute the forward difference

```

1 # Convert image_array to a data type that can handle negative numbers
2 image_array = image_array.astype('int16')
3
4 # Initialize gradient images with zeros for Forward Gradient
5 rows, cols = image_array.shape
6 forward_gradient_x = np.zeros_like(image_array, dtype='int16')
7 forward_gradient_y = np.zeros_like(image_array, dtype='int16')
8
9 # Compute Forward Gradient
10 for i in range(rows):
11     for j in range(cols - 1):

```

```

12     forward_gradient_x[i, j] = image_array[i, j + 1] - image_array[i, j]
13 for i in range(rows - 1):
14     for j in range(cols):
15         forward_gradient_y[i, j] = image_array[i + 1, j] - image_array[i, j]

```

---

## 1.2 Backward Difference

As mentioned above we now develop the python code to compute the backward difference

```

1 # Initialize gradient images with zeros for Backward Gradient
2 backward_gradient_x = np.zeros_like(image_array, dtype='int16')
3 backward_gradient_y = np.zeros_like(image_array, dtype='int16')
4
5 # Compute Backward Gradient
6 for i in range(rows):
7     for j in range(1, cols):
8         backward_gradient_x[i, j] = image_array[i, j] - image_array[i, j - 1]
9 for i in range(1, rows):
10    for j in range(cols):
11        backward_gradient_y[i, j] = image_array[i, j] - image_array[i - 1, j]

```

---

## 1.3 Central Difference

As mentioned above we now develop the python code to compute the central difference

```

1 # Initialize gradient images with zeros for Central Gradient
2 central_gradient_x = np.zeros_like(image_array, dtype='int16')
3 central_gradient_y = np.zeros_like(image_array, dtype='int16')
4
5 # Compute Central Gradient
6 for i in range(1, rows - 1):
7     for j in range(1, cols - 1):
8         central_gradient_x[i, j] = (image_array[i, j + 1] - image_array[i, j - 1]) // 2
9         central_gradient_y[i, j] = (image_array[i + 1, j] - image_array[i - 1, j]) // 2

```

---

You can see that we have used `dtype = 'int16'` that is to convert the data type of the image array to signed 16-bit integers. This is done to handle negative values that can occur when calculating gradients using differences between pixel intensities.

In the original image array, the data type is usually uint8, which supports unsigned integers ranging from 0 to 255. When calculating gradients, the difference between adjacent pixel values can be negative, and uint8 cannot represent negative numbers. This leads to overflow issues. Switching to int16 allows the array to store negative numbers, thus avoiding any overflow and accurately representing the gradients.

## 1.4 Results

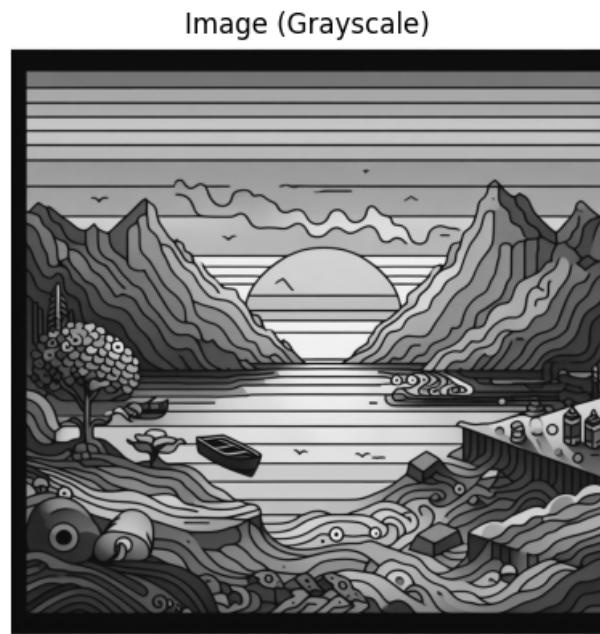


Figure 1: Orginal Image

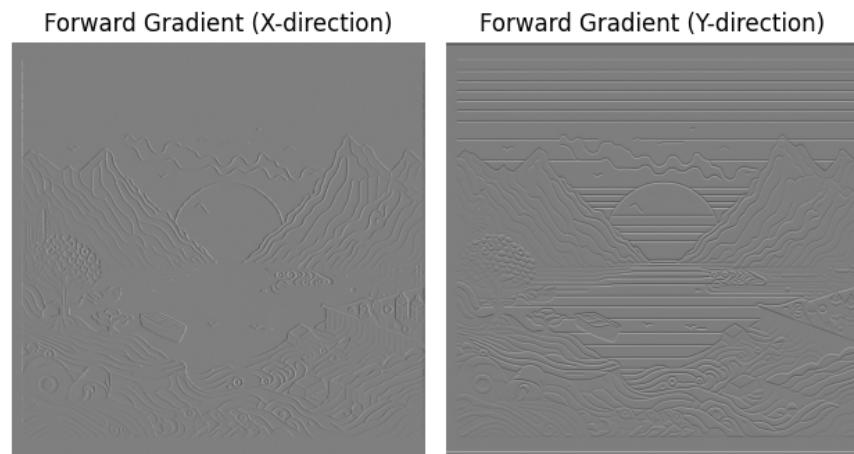


Figure 2: Forward Gradient in X and Y Direction

Backward Gradient (X-direction)      Backward Gradient (Y-direction)

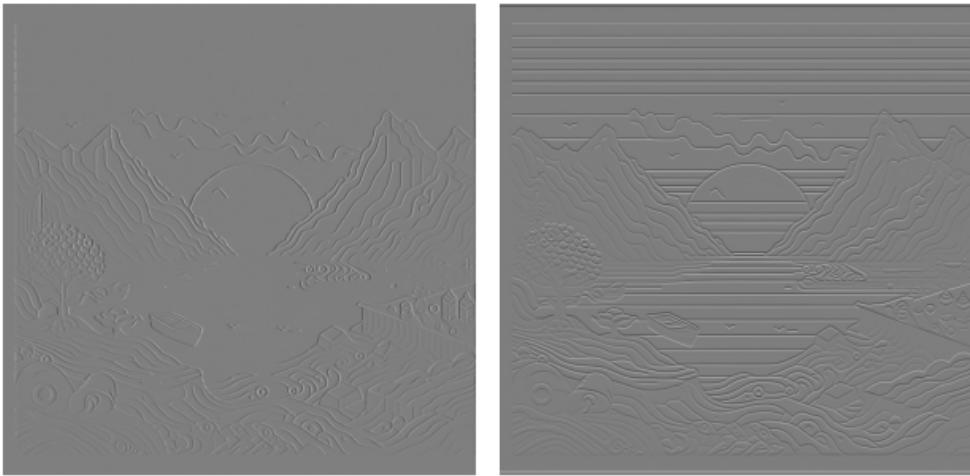


Figure 3: Backward Gradient in X and Y Direction

Central Gradient (X-direction)      Central Gradient (Y-direction)

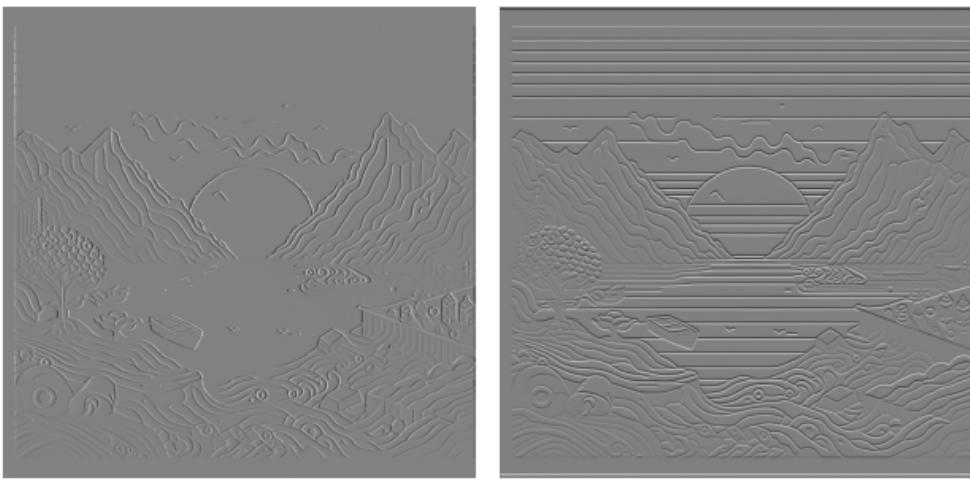


Figure 4: Central Gradient in X and Y Direction

The gradient magnitude  $M$  and direction  $\theta$  are calculated as:

$$M(i, j) = \sqrt{G_x^2(i, j) + G_y^2(i, j)}$$

$$\theta(i, j) = \arctan 2(G_y(i, j), G_x(i, j))$$

Normally taking the  $M(i, j) = \sqrt{G_x^2(i, j) + G_y^2(i, j)}$  is showing an warning that is due to overflow in the data type being used. The int16 data type can handle values from -32,768 to 32,767. When you square the values, they can easily exceed this range.

To avoid this, we can convert the gradient arrays to a floating-point data type (like float32 or float64) before performing the squaring and square root operations. This will prevent overflow issues.

## Farward

---

```
1 # Convert the gradients to float32 to prevent overflow
2 forward_gradient_x_float = forward_gradient_x.astype('float32')
3 forward_gradient_y_float = forward_gradient_y.astype('float32')
4
5 # Compute the gradient magnitude
6 gradient_magnitude_forward = np.sqrt(forward_gradient_x_float**2 +
    forward_gradient_y_float**2)
7
8 # Compute the gradient direction
9 gradient_direction_forward = np.arctan2(forward_gradient_y_float,
    forward_gradient_x_float) * 180 / np.pi
10
11 gradient_magnitude_forward, gradient_direction_forward
```

---

## Backward

---

```
1 # Convert the gradients to float32 to prevent overflow
2 backward_gradient_x_float = backward_gradient_x.astype('float32')
3 backward_gradient_y_float = backward_gradient_y.astype('float32')
4
5 # Compute the gradient magnitude
6 gradient_magnitude_backward = np.sqrt(backward_gradient_x_float**2 +
    backward_gradient_y_float**2)
7
8 # Compute the gradient direction
9 gradient_direction_backward = np.arctan2(backward_gradient_y_float,
    backward_gradient_x_float) * 180 / np.pi
10
11 gradient_magnitude_backward, gradient_direction_forward
12 gradient_magnitude_central, gradient_direction_central
```

---

## Central

---

```
1 # Convert the gradients to float32 to prevent overflow
2 central_gradient_x_float = central_gradient_x.astype('float32')
3 central_gradient_y_float = central_gradient_y.astype('float32')
4
5 # Compute the gradient magnitude
6 gradient_magnitude_central = np.sqrt(central_gradient_x_float**2 +
    central_gradient_y_float**2)
7
8 # Compute the gradient direction
9 gradient_direction_central = np.arctan2(central_gradient_y_float,
    central_gradient_x_float) * 180 / np.pi
10
11 gradient_magnitude_central, gradient_direction_central
```

---

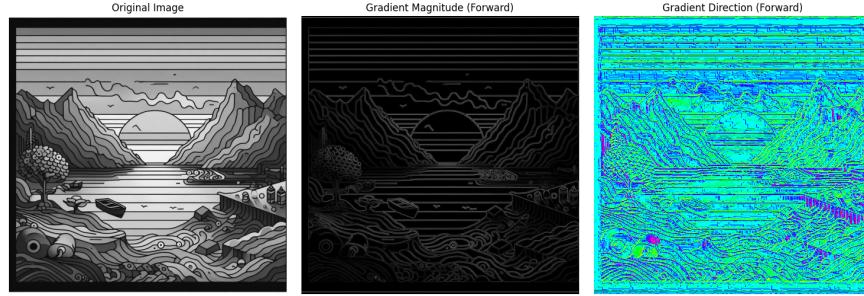


Figure 5: Gradient Magnitude and Gradient Direction in Forward Difference

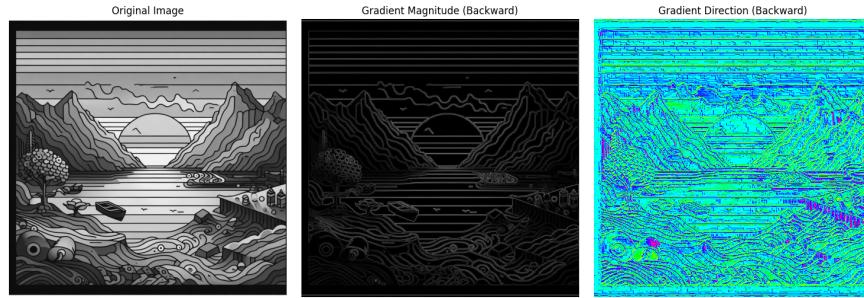


Figure 6: Gradient Magnitude and Gradient Direction in Backward Difference

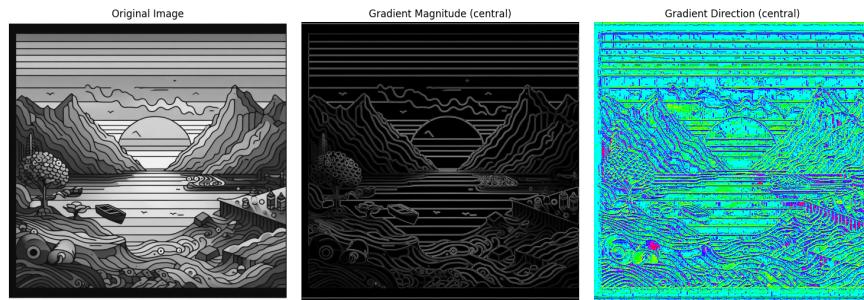


Figure 7: Gradient Magnitude and Gradient Direction in Central Difference

## 1.5 Anaylsis

**Gradient Magnitude Image :** In the gradient magnitude image we can see the edges or transition regions, the brighter regions correspond to higher magnitude which is directly proportional to the intensity of edges. where as the darker regions have indicate little to no change in intensity values, meaning uniform regions in the original image. The magnitude image is useful for identifying areas of interest in the original image, such as boundaries, outlines, or significant features.

**Gradient Direction Image :** We can see that the colors in the gradient direction image represent the orientation of the gradient at each pixel. This is an indication of the local orientation of edges or any features. some areas might appear to have random or mixed hues, indicating regions where the edge direction varies rapidly or where there is noise above in the images we can see such random hues, and the direction image is valuable for understanding the structure and orientation of the features within the image.

## 2 Question 2

Compute the LoG, DoG, and Gaussian pyramid scale space of the image

### Answer 2

Here we will be using the same image which we used for the above question to compute the Laplacian of Gaussian , Difference of Gaussian and Gaussian pyramid scale space of the image.

#### 2.1 Laplacian of Gaussian (LoG) :

The Laplacian of Gaussian (LoG) is a combination of two separate processes:

1. Gaussian Smoothing: This involves convolving the image with a Gaussian filter.
2. Laplacian Filtering: This is the application of a Laplacian filter to the image, which highlights regions of rapid intensity change.

The equation for the LoG for a 2D function is given by:

$$LoG(x, y, \sigma) = -\frac{1}{\sqrt{2\pi}\sigma^3} \left( 2 - \frac{x^2 + y^2}{\sigma^2} \right) e^{-\frac{x^2+y^2}{2\sigma^2}}$$

The liberty is given to use the builtin functions for LOG and DOG with the understanding of the function.

---

```
1 import cv2
2 # Gaussian smoothing
3 sigma = 1.0
4 gaussian_blur = cv2.GaussianBlur(image_array, (0, 0), sigma)
5 # Laplacian filtering
6 laplacian = cv2.Laplacian(gaussian_blur, cv2.CV_64F)
```

---

1.  $\sigma = 1.0$  : This line defines the standard deviation ( $\sigma$ ) of the Gaussian filter. The value of  $\sigma$  determines the extent of smoothing. A higher  $\sigma$  will lead to more blurring, while a lower  $\sigma$  will produce a sharper image.
2. `cv2.GaussianBlur(image_array, (0, 0), sigma)` : This is the main function that applies the Gaussian blur. The second argument,  $(0, 0)$ , defines the kernel size. By setting it to  $(0, 0)$ , you're allowing OpenCV to automatically determine the size based on the provided  $\sigma$  value.
3. The second argument in `cv2.Laplacian(gaussianblur, cv2.CV64F)` , `cv2.CV_64F`, specifies the depth of the resulting image. It indicates that the result should be a floating-point image. This is crucial because the Laplacian filter can produce both positive and negative variations, indicating dark-to-light and light-to-dark transitions, respectively

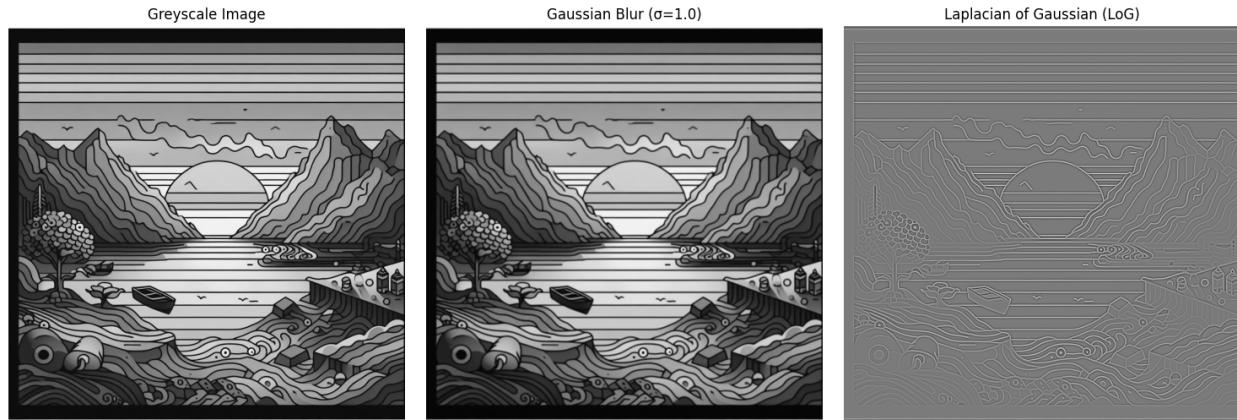


Figure 8: LoG

## 2.2 Difference of Gaussians (DoG)

The Difference of Gaussians (DoG) is closely related to the Laplacian of Gaussian (as we discussed above) we take Difference of 2 Gaussian Blurred, and is primarily used for blob detection.

1. Blur the image with two different Gaussian filters, where one has a slightly larger standard deviation  $\sigma$  than the other.
2. Subtract one blurred image from the other.

---

```

1 # Define two different sigmas
2 sigma1 = 1.0
3 sigma2 = 1.5
4
5 # Apply Gaussian blurring with the two sigmas
6 gaussian.blur1 = cv2.GaussianBlur(image_array, (0, 0), sigma1)
7 gaussian.blur2 = cv2.GaussianBlur(image_array, (0, 0), sigma2)
8
9 # Compute the DoG
10 dog = gaussian.blur2 - gaussian.blur1

```

---

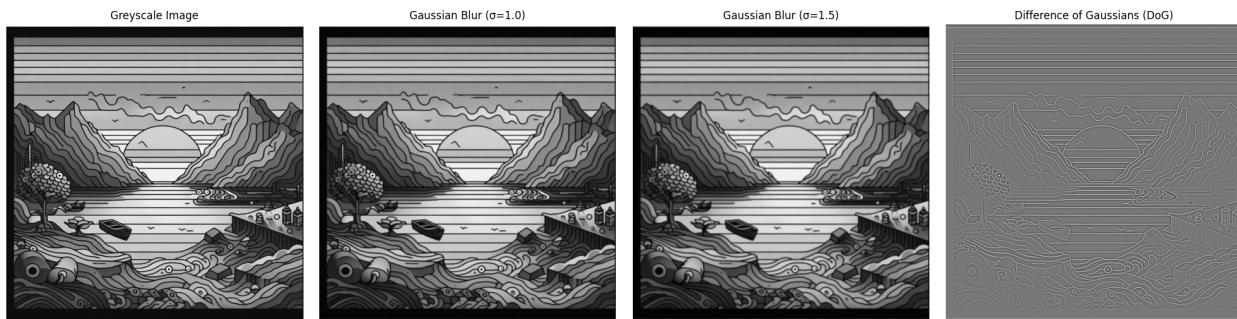


Figure 9: DoG

## 2.3 Gaussian Pyramid

The Gaussian pyramid is a multi-scale representation of an image. It is constructed by repeatedly reducing the size of the image and smoothing it with a Gaussian filter at each level. Each level of the pyramid is called a "scale".

1. Start with the original image.
2. Smooth the image using a Gaussian filter.
3. Reduce the size of the smoothed image by half in each dimension (typically, by sub-sampling every alternate pixel in both x and y directions).
4. Repeat the above two steps for the reduced image to get the next level of the pyramid.

We define the needed functions such as a function to convolve the gaussian filter to the image , a function to define the generate the guassian filter according to the size and the  $\sigma$  value and a function to genrate the different levels of guassian pyramid by using above 2 functions and down-sampling the image.

1.

---

```
1 def generate_gaussian_filter(size, sigma):
2     # Initialize the kernel with zeros
3     kernel = np.zeros((size, size))
4
5     # Calculate the center of the kernel
6     center = size // 2
7
8     # fill the kernel with values from the Gaussian function
9     for i in range(size):
10         for j in range(size):
11             x = i - center
12             y = j - center
13             kernel[i, j] = (1/(2*np.pi*sigma**2)) * np.exp(-(x**2 + y**2) / (2*sigma**2))
14
15     # Normalize the kernel
16     kernel /= kernel.sum()
17
18     return kernel
```

---

2.

---

```
1 def pad_image(image, pad_size):
2     return np.pad(image, ((pad_size, pad_size), (pad_size, pad_size)), mode='constant')
3
4 def convolve2D(image, kernel):
5     m, n = kernel.shape
6     pad_size = m // 2  # Assuming kernel is always odd-sized
7     padded_image = pad_image(image, pad_size)
8
9     y, x = padded_image.shape
10    y = y - m + 1
11    x = x - n + 1
12    new_image = np.zeros((y,x))
```

---

```

13     for i in range(y):
14         for j in range(x):
15             new_image[i][j] = np.sum(padded_image[i:i+m, j:j+n] * kernel)
16     return new_image

```

---

3.

---

```

1 def gaussian_pyramid(image, num_levels, sigma):
2     pyramid = [image]
3
4     for i in range(num_levels - 1):
5         # Create a Gaussian filter
6         gaussian_filter = generate_gaussian_filter(5, sigma)
7
8         # Convolve the image using the Gaussian filter
9         smoothed_image = convolve2D(pyramid[-1], gaussian_filter)
10
11        # Downsample the smoothed image
12        reduced_image = downsample(smoothed_image)
13
14        # Append the reduced image to the pyramid
15        pyramid.append(reduced_image)
16
17    return pyramid

```

---

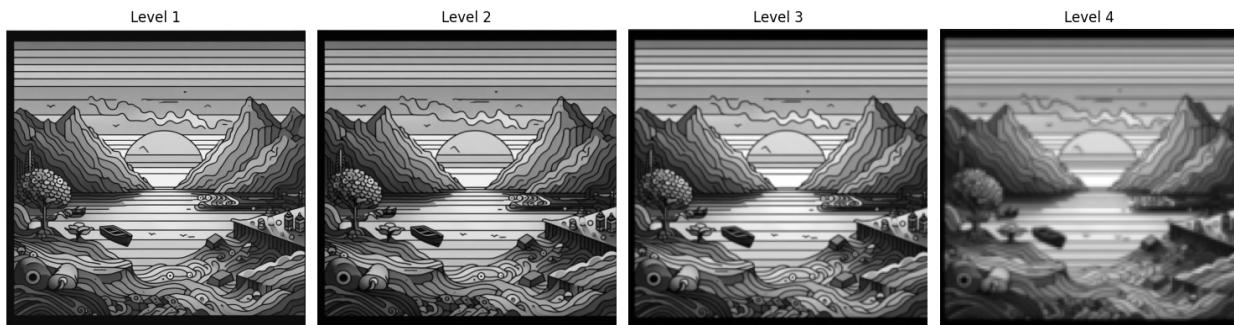


Figure 10: Gaussian Pyramid

---

### 3 Question 3

Develop the algorithm to merge the apple and orange using the pyramids as discussed in the class

### Answer 3

#### Algorithm to merge apple and orange using pyramids

1. Generate Laplacian pyramid of orange image.
2. Generate Laplacian pyramid of apple image.

3. Generate Combined Laplacian pyramid by copying left half of nodes at each level from apple

and right half of nodes from orange pyramids.

4. Reconstruct combined image from Combined Laplacian pyramid.

Already we have the code for getting the Gaussian pyramid of an image as discussed in previous question now we will build upon that to get the Laplacian pyramid of an image.

1. We get the Gaussian pyramid of the image and then proceed like this to get the Laplacian pyramid, we need to up sample the Gaussian pyramid level and subtract it with the before one.

---

```
1 # upsampling the neighbouring level
2 def upsample2x(image):
3     return image.repeat(2, axis=0).repeat(2, axis=1)
4
5 # Laplacian Pyramid
6 def laplacian_pyramid(gaussian_pyramid):
7     laplacian_pyramid = []
8     for i in range(len(gaussian_pyramid) - 1):
9         upsampled = upsample2x(gaussian_pyramid[i + 1])
10        if upsampled.shape != gaussian_pyramid[i].shape:
11            upsampled = upsampled[:gaussian_pyramid[i].shape[0], :gaussian_pyramid[i].
12                shape[1]]
13        laplacian = gaussian_pyramid[i] - upsampled
14        laplacian_pyramid.append(laplacian)
15    laplacian_pyramid.append(gaussian_pyramid[-1])
16    return laplacian_pyramid
```

---

2. Now after getting the Laplacian pyramid we do blending and reconstruct the image as mentioned above.

---

```
1 def blend_laplacian_pyramids(laplacian1, laplacian2):
2     blended_pyramid = []
3     for i in range(len(laplacian1)):
4         # Take the left half from the apple image and the right half from the orange
5         # image
6         rows, cols = laplacian1[i].shape
7         blended = np.hstack((laplacian1[i][:, :cols//2], laplacian2[i][:, cols//2:]))
8         blended_pyramid.append(blended)
9     return blended_pyramid
10 # Function to reconstruct an image from a Laplacian pyramid
11 def reconstruct_from_laplacian(laplacian_pyramid):
12     reconstructed_image = laplacian_pyramid[-1] # Start with the smallest image
13     for i in range(len(laplacian_pyramid) - 2, -1, -1):
14         # Upsample and add to the next level
15         upsampled = upsample2x(reconstructed_image)
16         if upsampled.shape != laplacian_pyramid[i].shape:
17             upsampled = upsampled[:laplacian_pyramid[i].shape[0], :laplacian_pyramid[i].
18                 shape[1]]
19         reconstructed_image = upsampled + laplacian_pyramid[i]
20     return reconstructed_image
```

---

**Blending Laplacian Pyramids Function :** The *blend\_laplacian\_pyramids* function combines two Laplacian pyramids by horizontally stacking the left half of each image from the first pyramid with the right half from the second pyramid. This results in a new pyramid where each level has the left side sourced from the first image and the right side from the second image, effectively blending the two original images.

**Reconstructing from Laplacian Pyramid Function :** The *reconstruct\_from\_laplacian* function rebuilds an image from its Laplacian pyramid representation. Starting with the smallest image in the pyramid, it successively upsamples and adds the current image to the next pyramid level. This process reconstructs the original (or in the case of our blend, the combined) image from its multi-resolution representation in the Laplacian pyramid.

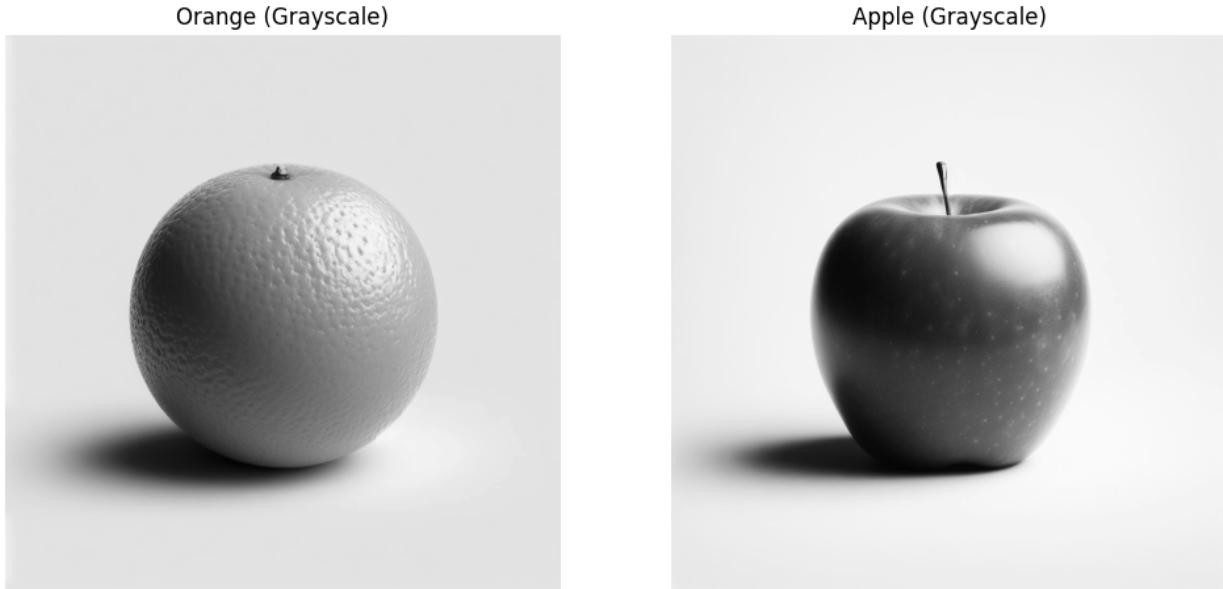


Figure 11: Orange and Apple Images

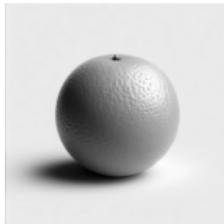
Orange Level 0



Apple Level 0



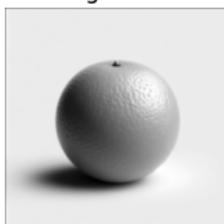
Orange Level 1



Apple Level 1



Orange Level 2



Apple Level 2



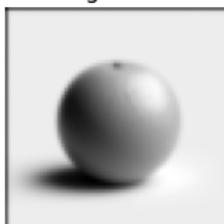
Orange Level 3



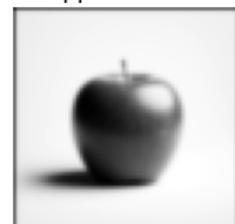
Apple Level 3



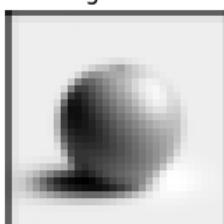
Orange Level 4



Apple Level 4



Orange Level 5



Apple Level 5

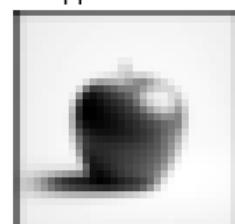


Figure 12: Gaussian Pyramid

Orange Laplacian Level 0



Apple Laplacian Level 0



Orange Laplacian Level 1



Apple Laplacian Level 1



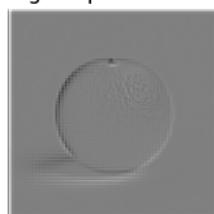
Orange Laplacian Level 2



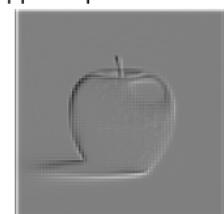
Apple Laplacian Level 2



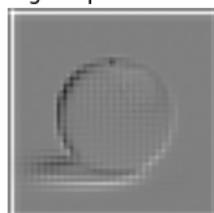
Orange Laplacian Level 3



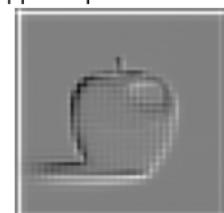
Apple Laplacian Level 3



Orange Laplacian Level 4



Apple Laplacian Level 4



Orange Laplacian Level 5



Apple Laplacian Level 5

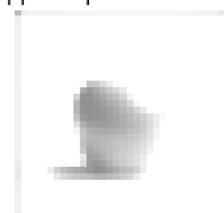


Figure 13: Laplacian Pyramid

## Blended Image

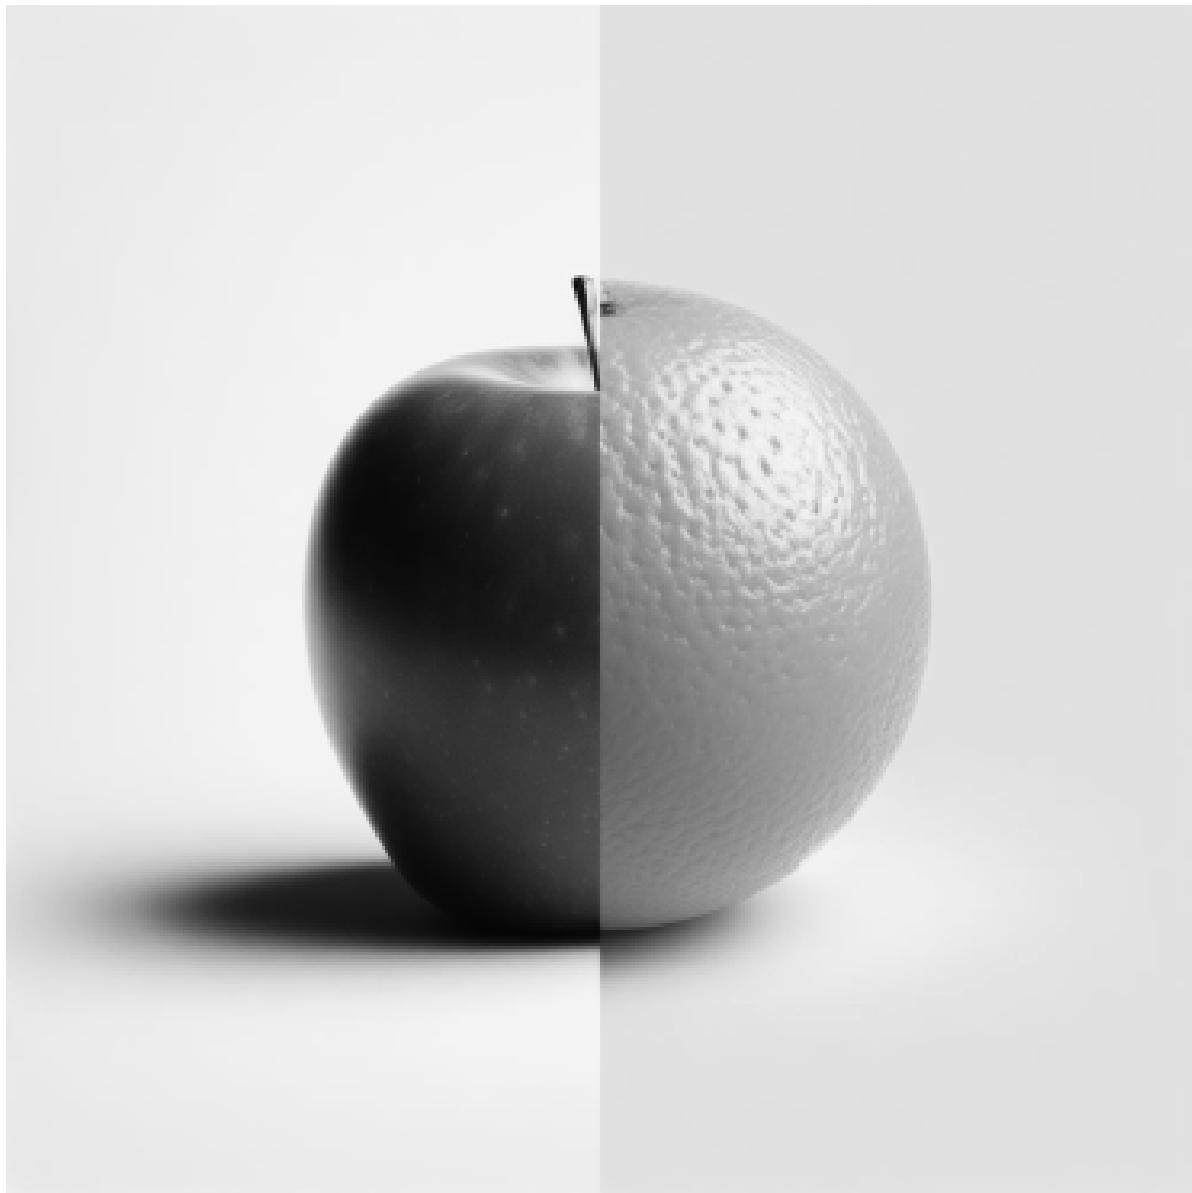


Figure 14: Combined Image